

```
In [ ]: import numpy as np
import pandas as pd
import io #google colab
from google.colab import files #google colab
import scipy.stats as stats
import statsmodels.stats.weightstats as stest
from statsmodels.stats.weightstats import ztest
from statsmodels.stats.proportion import proportions_ztest
from statsmodels.stats.power import ttest_power
from scipy.stats import ttest_1samp, ttest_ind, ttest_rel, kruskal #kruskal is replacement of annova for non parametric test
import statsmodels.api as sm
from statsmodels.formula.api import ols #ols=ordinaryLeastSquares. Creates a model from dataframe
from statsmodels.stats.anova import anova_lm
from statsmodels.stats.multicomp import pairwise_tukeyhsd
import matplotlib.pyplot as plt
import seaborn as sns
from math import factorial
```

/usr/local/lib/python3.7/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.

```
import pandas.util.testing as tm
```

```
In [ ]: #Binomial (discrete)
# for finding probability of exact k true outcomes out of n total outcomes where probability of true is p
scipy.stats.binomial.pmf(k,n,p)

#other ones

scipy.stats.binomial.cdf()#cumulative probability
scipy.stats.binomial.sf()#survival function
```

```
In [ ]: #uniform distribution (discrete)
#probability of choosing k as value of x out of range min value of x as low and max value of x as high(max+1) when all values of x are equally likely
randint.pmf(k,lowm,high)

#probability of choosing 3 or 4 out of {1,2,3,4,5,6}
req_prob=randint.pmf(k=3,low=1,high=7)+randint.pmf(k=4,low=1,high=7)

#uniform distribution (continuous)

scipy.stats.uniform.pmf()
scipy.stats.uniform.cdf()
scipy.stats.uniform.sf()

#for a range 0-200 of continuous uniform distribution
#probability of 50-150 can be calculated by cum_prob_150 - cum_prob_50
req_prob=scipy.stats.uniform.cdf(x=150,loc=0,scale=200)-scipy.stats.uniform.cdf(x=50,loc=0,scale=200)
```

```
In [ ]: #Normal(continuous)
scipy.stats.normal.pdf()#single range #the value of pdf at a particular point is always zero 0

#when required value is a, avg=x, var=y std=np.sqrt(y) z=(a-x)/std

scipy.stats.normal.cdf(z)#cumulative(left), probability of value less than a

scipy.stats.normal.sf(z) #(1-cdf) cumulative right, probability of value greater than a
```

```
In [ ]: # sampling
random.choices() #draw sample with replacement
random.sample() #draw sample without replacement
```

```
In [ ]: #bernoulli is used when only two outcomes are possible (head/tail)(goal/miss) (pass/fail)
#calculate probability for a goal for a free kick(x=k=1) for a player with 70% (p=0.7) goal rate
scipy.stats.bernoulli.pmf(k=1,p=0.7)
scipy.stats.bernoulli.cdf()
```

```
In [ ]: #poission
scipy.stats.poisson.pmf()
scipy.stats.poisson.cdf()
scipy.stats.poisson.sf()
```

```
In [ ]: #if x is desired outcome of n samples which belong to population N
#the number of desired outcomes out of Population N can be given by:
N=50
n=13
x=8
p_samp=x/n
#Estimation of desired outcome from the population
num_out=N*p_samp
```

misc

Coefficient of variance= std/mean..... is used for selecting a feature with least variance

Margin Of error/sampling error = $z * \text{std} / \sqrt{n}$

sampling error=samplemean-pop.mean.....when no confidence interval available

```
In [ ]: #margin of error

# Margin Of error/sampling error = z * std / sqrt(n)

#sample n
n=100
std=0.15 #=15%
conf_level=90%
#which makes alpha = 1-conf_interval = 10% =0.1
#since we are trying to find interval q= alpha/2
z_alpha_by_2=np.abs(round(stats.norm.isf(q=0.1/2),4))#isf gives z stat value
error=(z_alpha_by_2 * std)/np.sqrt(n)
```

```
In [ ]: #confidence interval when mean and std are given(continuous)

x_bar=45000 #sample_avg
conf_interval=0.95 #95% conf interval
pop_std=4000
n=250
interval_95=stats.norm.interval(conf_interval,loc=x_bar,scale=pop_std/np.sqrt(n))
print(interval_95)

interval_99=stats.norm.interval(alpha=0.99,loc=x_bar,scale=pop_std/np.sqrt(n))
print(interval_99)

(44504.16397415635, 45495.83602584365)
(44348.361002958416, 45651.638997041584)
```

```
In [ ]: #confidence interval when probability is given(discrete)  
N=50  
n=13  
x=8  
p=8/13  
alpha=0.99  
stats.norm.interval(0.99,loc=p,scale=np.sqrt((p*(1-p))/n))
```

```
Out[ ]: (0.26782280814713805, 0.9629464226220927)
```

The point estimate for the average is `np.mean()` of the sample

Check for normality Shapiro

H0: data is normally distributed(skew=0)

H1: data is not normally distributed(skew!=0)

Check for variance Levene

H0: variance are equal

H1: variance are not equal

```
In [ ]: stats.shapiro(x)#h1 variance are different  
stats.levene(sample1,sample2...samplen)#h0 variance are same
```

```

In [ ]: #hypothesis testing

#rejecting  $H_0$ :

#based on pval:
#pval<=significance(0.05)

#based on conf interval
# $\mu_0$  doesnt lie in confidence interval

##based on critical regions

#for two tailed test
# $\mu \neq \mu_0$ 
# $|z| > z_{\alpha/2}$ 

#for left tailed test
# $\mu < \mu_0$ 
# $z < -z_{\alpha}$ 

#for right tailed test
# $\mu > \mu_0$ 
# $z > z_{\alpha}$ 

# z_calc

z_calc=(x_bar- $\mu$ )/(pop_std/np.sqrt(n))

def z_test(pop_mean,pop_std,n,samp_mean):
    z_score=(samp_mean-pop_mean)/(pop_std/np.sqrt(n))
    return z_score

#z from pval
stats.norm.ppf(1-pval)

#pval
pval=stats.norm.cdf(z_score)#left tailed
pval=stats.norm.sf(z_score)#right tailed

#only for two tailed test
req_pval=pval*2

#interval
stats.norm.interval(alpha,loc=x_bar,scale=pop_std/np.sqrt(n)) #check if the cl
aimed value lies within the interval

#type1 error is probability of wrongly rejecting true null hypothesis(alpha)
#type2 error is porbability of wrongly accepting false null hypothesis(beta)
#(1-beta) is power of test i.e accpeting correctly rejecting false null hypoth
esis

```

```

#calculate power of ztest
power.zt_ind_solve_power() # from statsmodels library
#alpha is fixed beta is controlled

#ztest inbuilt:
statsmodels.stats.weightstats.ztest(x1=array,value=Mu,alternative='smaller'/'larger'/'two-sided')

#for less than 30 data points
stats.ttest_1samp(a=array,popmean=Mu)

#z_crit
stats.norm.isf(q)#q=1-alpha

```

Two sample Ztest

```

def TwoSampZTest(samp_mean_1, samp_mean_2, samp_std_1, samp_std_2, value, n1, n2):

denominator = np.sqrt((samp_std_1**2 / n1) + (samp_std_2**2 / n2))

zscore = ((samp_mean_1 - samp_mean_2) - (value)) / denominator

return zscore

```

```

In [ ]: #two sample ttest

#stats.ttest_ind(a=array1,b=array2)

#stats.ttest_rel(a=array1,b=array2)

```

```

In [ ]: z_stat=(75-60)/(30/np.sqrt(45))
print(z_stat)

print(stats.norm.sf(z_stat))

```

```

3.3541019662496843
0.0003981150787954058

```

z = stats.norm.isf(q=0.01/2)for z critical two sided

z = stats.norm.isf(q=0.010)for z critical one sided

Critical values:

stats.norm.isf(q=0.05)..... zdistribution

stats.t.isf(q=0.05,df=n-1).....tdistribution.....n is size of sample.....one sample

stats.t.isf(q=0.05,df=(n1+n2)-2).....tdistribution.....n is size of sample.....twosample

stats.chi2.isf(q=0.05,df=(cols-1)*(rows-1)).....if only one column is present its rows - 1

```
In [ ]: #test of proportion
#one sample
z = (sample_prop - specified_prop) / np.sqrt ( specified_prop * (1 - specified_
prop) / n)
#rejecting H0:

#based on pval:
#pval<=significance(0.05)

#based on conf interval
#Mu0 doesnt lie in confidence interval

##based on critical regions

#for two tailed test
#Mu!=Mu0
#abs(z)>z_alpha/2

#for left tailed test
#Mu<Mu0
#z<-z_alpha

#for right tailed test
#Mu>Mu0
#z>z_alpha

statsmodels.stats.proportion.proportions_ztest(count, nobs, value=test_prop, a
lternative='two-sided', prop_var=False)#for two sample value= 0

#standard error for proportion
se=z_crit*(np.sqrt(p*(1-p)/n))#
```

```
In [ ]: samp_prop=105/361
specified_prop=0.25
z=(samp_prop-specified_prop)/np.sqrt(specified_prop*(1-specified_prop)/361)
print(z)
pval=stats.norm.sf(z)
print(pval)
```

1.7928245201151534
0.03650049373124949

```

In [ ]: #Test of proportions
#two sample
p_bar=(x1+x2)/(n1+n2)=((n1*p1)+(n2*p2))/(n1+n2)#n1,n2 are samples p1,p2 are proportions of specific items

z=(p1-p2)/np.sqrt(p_bar*(1-p_bar)*((1/n1)+(1/n2)))
#based on pval:
#pval<=significance(0.05)

#based on conf interval
#Mu0 doesnt lie in confidence interval

##based on critical regions

#for two tailed test
#Mu!=Mu0
#abs(z)>z_alpha/2

#for left tailed test
#Mu<Mu0
#z<-z_alpha

#for right tailed test
#Mu>Mu0
#z>z_alpha

statsmodels.api.stats.proportions_ztest(Sample_1,Sample_2)

```

```

In [ ]: p1=78/100
p2=65/70
print(p1,p2)
p_bar=round(((100*p1)+(70*p2))/(100+70),2)
print(p_bar)
z=(p1-p2)/np.sqrt(p_bar*(1-p_bar)*((1/100)+(1/70)))
print(z)
p_val=stats.norm.cdf(z)
print(p_val)

```

```

0.78 0.9285714285714286
0.84
-2.600520156052018
0.00465412753578429

```



```

In [ ]: #non parametric test for one sample

scipy.stats.wilcoxon(diff,alternative)

sample=[9.1,10.1,9.9,9.9,10,9.8,9.7,9.8,9.9,8.6]
diff=[i-10 for i in sample]
scipy.stats.wilcoxon(diff,alternative='less')

#non parametric test for two sample paired

scipy.stats.wilcoxon(sample_1,sample_2,alternative)

before=[2.7,1.1,1.4,0.9,0.9]
after=[1.3,1.4,1.1,1.3,1.9]
scipy.stats.wilcoxon(before,after,alternative='greater')

#non parametric test for two sample unpaired

##ranksums
scipy.stats.stats.ranksums(sample1,sample2)

a=[32,31,33,46]
b=[49,20,58,55,52]
scipy.stats.stats.ranksums(a,b)#scipy.stats.stats.ranksums(a,b,alternative='less') may work if new version

#mannwhitneyu
scipy.stats.stats.mannwhitneyu(sample1,sample2,alternative)
a=[32,31,33,46]
b=[49,20,58,55,52]
scipy.stats.stats.mannwhitneyu(a,b,alternative='less')

```

```

In [ ]: #Non-Parametric equivalent of annova
scipy.stats.kruskal(sample1,sample2,sample3)

#Non-parametric posthoc
scikit_posthocs.posthoc_conover(dataframe,val_col,group_col)

```

```

In [ ]: #Chi-Square goodness of fit

#H0: obs=exp

#H1: obs!=exp

stats.chisquare(f_obs=observed vals,f_exp=expected vals)

```

```
In [ ]: #Chi-Square Test for Independence

#H0 : The variables are independent

#H1 : The variables are not independent (i.e. variables are dependent)

chistat,pval,dof,arr=stats.chi2_contingency(observed=pd.crosstab(df['CAT'],df
['CAT'],correction=False)
```

anova

It is used to check the equality of population means for more than two independent samples. Each group is considered as a treatment. It assumes that the samples are taken from normally distributed populations. To check this assumption we can use the Shapiro-Wilk Test. Also, the population variances should be equal; this can be tested using the Levene's Test.

The null and alternative hypothesis is given as:

H0 : The averages of all treatments are the same.

H1 : At least one treatment has a different average.

```
In [ ]: #A=df[df['ethnicity']=='group A']['Total_score']

#B=df[df['ethnicity']=='group B']['Total_score']

#C=df[df['ethnicity']=='group C']['Total_score']

#perform shapiro and Levene on ABC

#if data is normal and having equal variance

stats.f_oneway(A,B,C)
```

aov_table

formula='avg_score ~ ethnicity' #one way annova

formula='avg_score ~ ethnicity+other categorical+ethnicity:other category'#two way annova

model=ols(formula,df).fit()

aov_table=anova_lm(model,typ=1)

aov_table

```
In [ ]: #post-hoc

#tukeyhsd when equal number of obsrvations are present else scheffe

MultiComp=statsmodels.sandbox.stats.multicomp.MultiComparison(data=df_machine
['Strength'],groups=df_machine['machine'])
MultiComp.tukeyhsd().summary()

#scheffe
```

```
Out[ ]: MannwhitneyuResult(statistic=4.0, pvalue=0.08895479749349122)
```

```
In [ ]: X= statsmodels.stats.multicomp.pairwise_tukeyhsd(endog=independent, groups=cat
egory, alpha= 1-CI)#aplha is 0.05 in case of 95% conf interval

>>> from statsmodels.examples.try_tukey_hsd import cylinders, cyl_labels
>>> from statsmodels.stats.multicomp import MultiComparison
>>> cardata = MultiComparison(cylinders, cyl_labels)
>>> results = cardata.tukeyhsd()
>>> results.plot_simultaneous()
```

NON parametric alternative for anova

H0: median is equal for all groups

H1: median is not equal for all groups

stats.kruskal(a-group,b-group,c-group)

Probability

combination

```
def com(n,r):
```

```
    comb = factorial(n)/(factorial(r)* factorial(n-r))
```

```
    return comb
```

```
prob = com(number of desired outcomes in set ,r) / com(all outcomes in set ,r)
```

alternative

```
from itertools import combinations
```

```
a=len(list(combinations([array of desired outcomes],r))
```

```
b=len(list(combinations([array of all outcomes],r))
```

```
p=a/b
```

Mutually exclusive

Mutually exclusive events arent independent

if events are mutually exclusive then addup the probabilities

addition rule for mutually exclusive events $p(A \cup B) = p(A) + p(B)$

addition for mutually non exclusive events $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

mutually independent

multiplication rule for independent events $p(A \cap B) = p(A) \cdot p(B)$

alternatively $p(A) = p(A|B)$ for mutually independent events

multiplication rule for non-independent events $p(A \cap B) = p(A) \cdot p(B|A)$

alternatively $p(B|A) = p(A \cap B) / p(A)$

Bayes theorem

$p(h|e) = (p(e|h) / p(e)) \cdot p(h)$

Binomial Probability distribution function discrete

$$P(x) = n! / (x! * (n-x)!)$$

$p(x)$ is probability of getting x successes in n trials

mean/avg of binomial distribution is $n * p$

$$\text{std} = \sqrt{np(1-p)}$$

n is number of trials p is probability of success.

`stats.binom.pmf(k=array,n=nobs,p=given prob)`.....output is probability of each element in array

`stats.binom.cdf(k=array,n=nobs,p=given prob)`.....for left hand side probability

`stats.binom.sf(k=array,n=nobs,p=given prob)`.....for right hand side probability

poisson Distribution % or rate

`stats.poisson.pmf(n=array of possible numbers, rate=given rate)`.....gives an array of probability of all the possible numbers.

if the array passed is in form of 0,1,2,3 `poisson [4]` gives the probability of number 4

for a probability of 3 or less than 3 is given by `poisson[0]+poisson[1]+poisson[2]+poisson[3]`

for a probability of greater than 3 is given by `1 - (poisson[0]+poisson[1]+poisson[2]+poisson[3])`

normal Probability distribution function/gaussian distribution continuous

mean=median=mode

+1 std has 68% of the data 2/3

+2 std has 95% of the data 19/20

+3 std has 99.7% of the data 997/1000

for problem a - Calculate the z value

$$z = (x - \mu) / \text{std}$$

`stats.norm.cdf(z)`.....this outputs probability for less than x

`stats.norm.cdf(x,loc=Mu,scale=std)`....optimised form using the original values

`1-stats.norm.sf(x,loc=Mu,scale=std)`...this outputs probability for greater than x

`stats.norm.cdf(x2,loc=Mu,scale=std)-stats.norm.cdf(x1,loc=Mu,scale=std)`....this outputs the probability of values between x_2 and x_1 where $x_2 > x_1$

for interval

<sample size greater than 30>

`stats.norm.interval(alpha=0.9,loc=xbar,scale=std/np.sqrt(n)).....`

<sample size less than 30>

`stats.t.interval(alpha=0.9,loc=xbar,scale=std/np.sqrt(n),df=n-1).....`

mathematical formula

`lower_value=round(np.mean(samp)-z*np.std(samp)/np.sqrt(n))`

`upper_value=round(np.mean(samp)+z*np.std(samp)/np.sqrt(n))`

proportion interval formula

`p_samp=true_obs/sample`

`stats.norm.interval(alpha,loc=p_samp,scale=np.sqrt(p_samp*(1-p_samp)/n)`

```
In [ ]: #stats.wilcoxon(array1,array2,alternatives=({"two-sided", "greater", "Less"})).....paired sample
        #stats.mannwhitneyu(array1,array2,alternatives=(None, 'two-sided', 'less', 'greater')).....independent sample when arrays are of same size
        #stats.ranksums(array1,array2)..... independent samples
        #stats.wilcoxon(x-checkVal,{"two-sided", "greater", "less"})
```

```
In [ ]:
```

Prince Distribution Plot one sided tests.

```

In [ ]: n = 75
pop_mean = 175
samp_mean = 176.3
samp_std = 8
CI = 0.95
alpha = 0.05 # use this for checking

# H0: mu<=175
# H1: mu>175
def z_test(pop_mean,samp_std,n,samp_mean):
    zscore=(samp_mean-pop_mean)/(samp_std/np.sqrt(n))
    return zscore
z_score=z_test(pop_mean,samp_std,n,samp_mean)
print(z_score)
#pval
pval=stats.norm.cdf(z_score)
print(pval)
## plotting.
plt.rcParams['figure.figsize'] = [15,8]

X = stats.norm.rvs(size = 20000, random_state = 10)
ax = sns.distplot(a = X, hist = False, kde_kws = dict(cumulative=False))

plt.ylim(0, 0.42)

z = stats.norm.isf(q = 0.05)
print('z value:',z)
# plot a vertical line at z-value
x, y = [z, z], [0, 0.3]

# drawing z line
plt.plot(x, y, color = 'black')

# coloring
l1 = ax.lines[0]
x1 = l1.get_xydata()[:,0]
y1 = l1.get_xydata()[:,1]
ax.fill_between(x1[x1 >= z], y1[x1 >= z], color = 'lightblue')
# labeling the regions
plt.text(x = -1.15, y = 0.18, s = 'Acceptance Region for H0 (1 - α)')
plt.annotate(s = 'α = 0.05\n(Rejection Region for H0)', xy = (2.5, 0.005), xyt
ext = (3.5, 0.035),
            arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

plt.annotate(s = 'z_crit = '+ str(round(z,2)), xy = (z, 0.13), xytext = (z+0.
8, 0.15),
            arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

# general annotation.

plt.xticks(ticks = [-6, -4, 0, 2, 4, 6, round(z_score,2)],
           labels = [-6, -4, 0, 2, 4, 6, 'z_test = '+str(round(z_score,2))])

plt.xlabel('Z', fontsize = 15)

```

```
plt.ylabel('Probability', fontsize = 15)
plt.title('Right-Tailed Test', fontsize = 15)

plt.show()
```

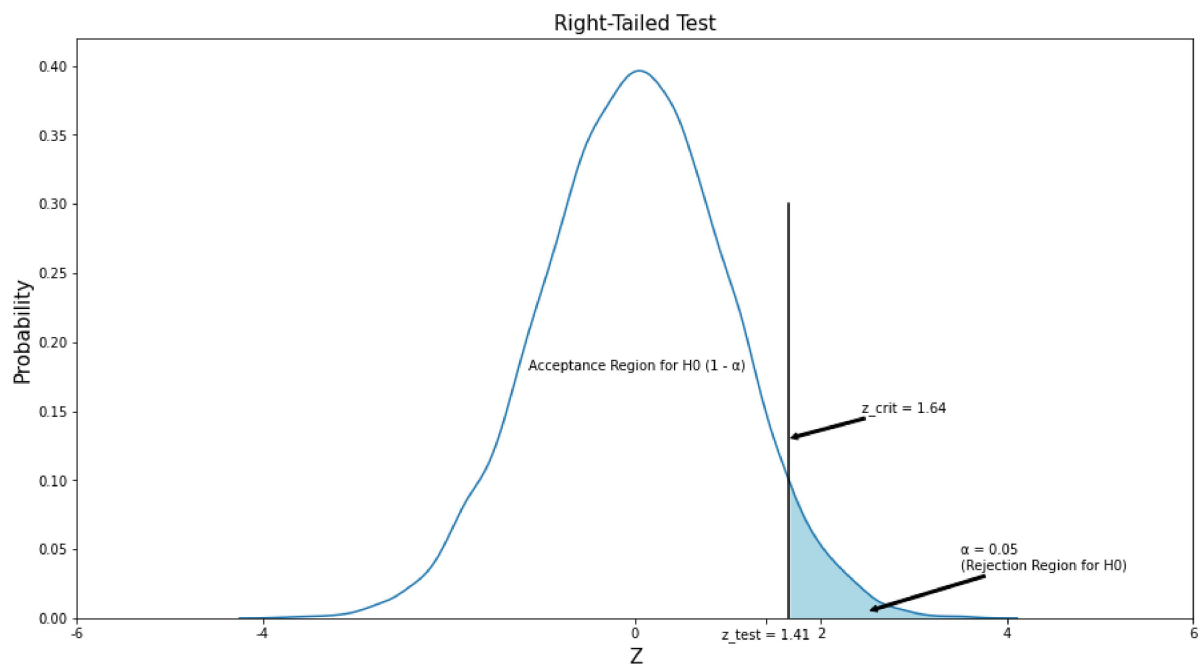
1.4072912811497253

0.920329483175748

z value: 1.6448536269514729

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)



Prince Distribution Plot Two sided tests


```

In [ ]: # type your code here
        # type your code here

# H0 : mu = 25
# H1 : mu != 25

n = 40
pop_mean = 25
pop_std = 4.8
samp_mean = 23.2

def z_test(pop_mean, pop_std, n, samp_mean):
    z_score = (samp_mean - pop_mean) / (pop_std / np.sqrt(n))
    return z_score

z_score = z_test(pop_mean, pop_std, n, samp_mean)
print(z_score)

CI = 0.95
z_critical = round(stats.norm.isf(q = 0.05),2)
print(z_critical)

plt.rcParams['figure.figsize'] = [15,8]

X = stats.norm.rvs(size = 20000, random_state = 10)
ax = sns.distplot(a = X, hist = False, kde_kws = dict(cumulative=False))

plt.ylim(0, 0.42)

z = stats.norm.isf(q = 0.05)
print('z value:',z)
# plot a vertical line at z-value

# +z value
x, y = [z, z], [0, 0.3]
plt.plot(x, y, color = 'black')
# -z value
x, y = [-z, -z], [0, 0.3]
plt.plot(x, y, color = 'black')

# coloring
l1 = ax.lines[0]
x1 = l1.get_xydata()[:,0]
y1 = l1.get_xydata()[:,1]
ax.fill_between(x1[x1 >= z], y1[x1 >= z], color = 'lightblue')
ax.fill_between(x1[x1 <= -z], y1[x1 <= -z], color = 'lightblue')

# labeling the regions
plt.text(x = -1.15, y = 0.18, s = 'Acceptance Region for H0 (1 - α)')
plt.annotate(s = 'α/2 = 0.05\n(Rejection Region for H0)', xy = (2, 0.018), xytext = (2.5, 0.05),
            arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

plt.annotate(s = 'α/2 = 0.05\n(Rejection Region for H0)', xy = (-2, 0.018), xytext = (-2.5, 0.05),
            arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

```

```

text = (-4.5, 0.055),
        arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

plt.annotate(s = 'z_crit = ' + str(round(z,2)), xy = (z, 0), xytext = (z-1.2,
0.02),
        arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

plt.annotate(s = 'z_crit = ' + str(round(-z,2)), xy = (-z, 0), xytext = (-z+0.
3, 0.02),
        arrowprops = dict(facecolor = 'black', arrowstyle = 'simple'))

# general annotation.

plt.xticks(ticks = [-6, -4, 0, 2, 4, 6, round(z_score,2)],
        labels = [-6, -4, 0, 2, 4, 6, 'z_test = ' + str(round(z_score,2))])

plt.xlabel('Z', fontsize = 15)
plt.ylabel('Probability', fontsize = 15)
plt.title('Right-Tailed Test', fontsize = 15)

plt.show()

```

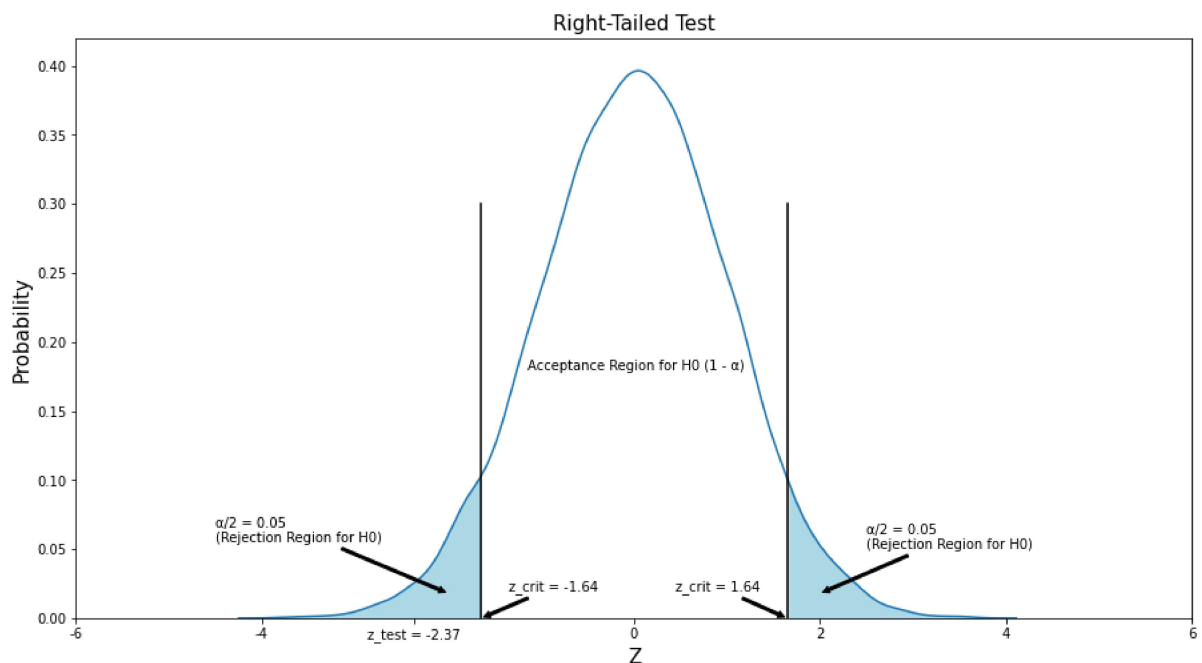
-2.3717082451262854

1.64

z value: 1.6448536269514729

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

warnings.warn(msg, FutureWarning)



In []: a